

!How causal mapping helps a real evaluation – the Love Alliance case

31 May 2026

By Steve and Gabriele at Causal Map, with Dena at Southern Hemisphere.

The Love Alliance is a five-year partnership, running from 2021 to 2025, that works to improve the health and rights of key populations affected by HIV, meaning sex workers, people who use drugs, LGBTQ+ people and people living with HIV, across ten countries in Africa and through regional and global advocacy. It is funded by the Dutch Ministry of Foreign Affairs, administered through Aidsfonds, and delivered by a consortium of partner organisations working through advocacy, capacity-strengthening and movement building. Its theory of change is a tidy ladder (Figure 1): strategies at the bottom, then short, medium and long-term outcomes, then goals at the top. Clear, communicable, fundable.

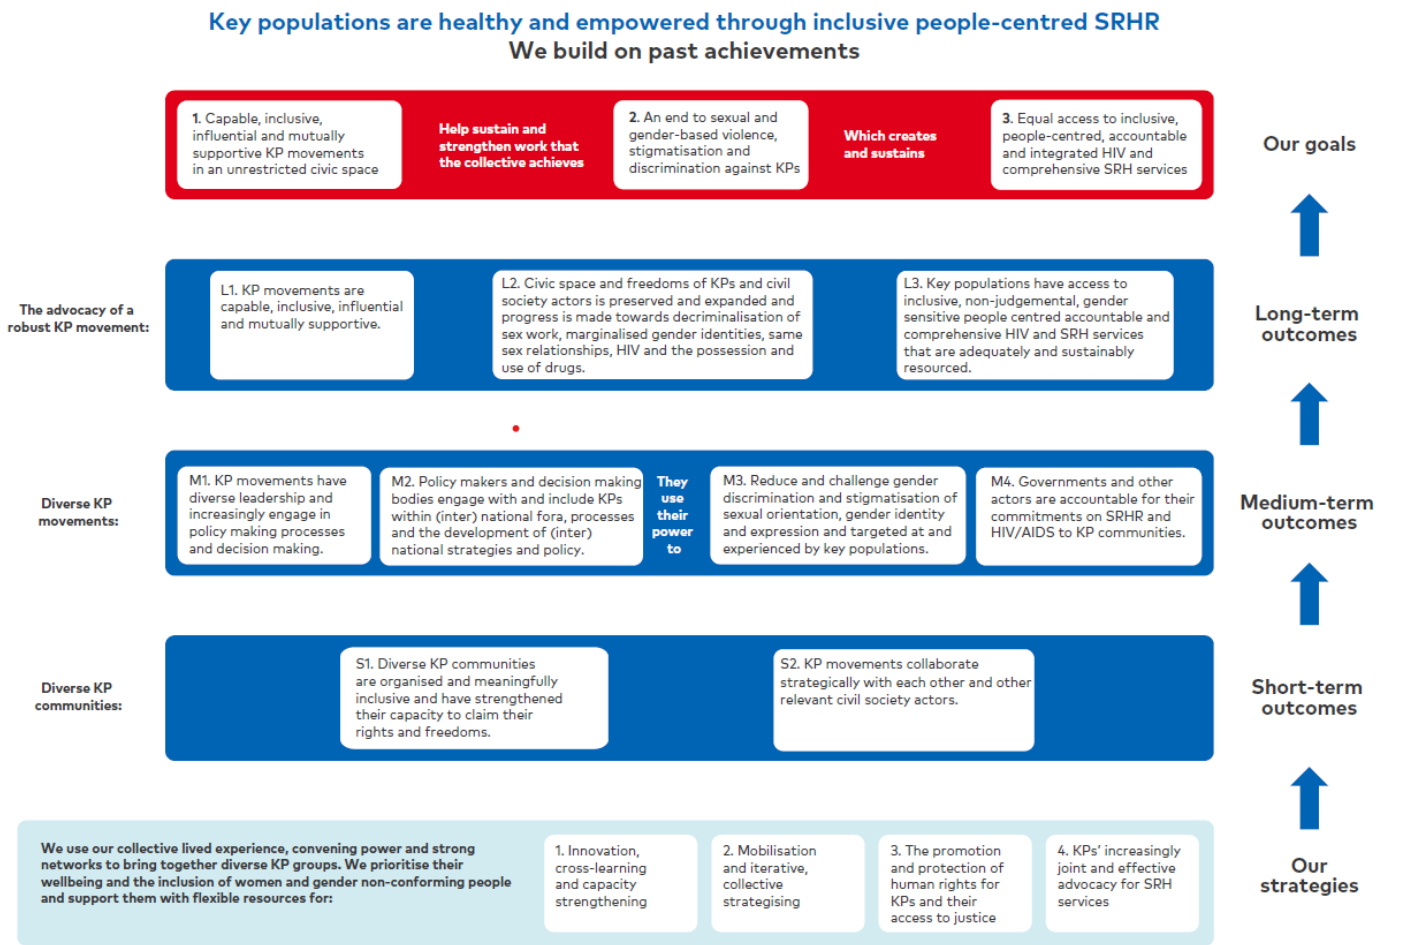


Figure 1. The programme's official theory of change: a layered ladder from strategies up to goals.

Southern Hemisphere, a research and evaluation consultancy, was asked to carry out the end-term evaluation of the whole partnership: how well it had worked, how plausibly it had contributed to change, and whether the gains could last. They did this with a participatory, mixed-methods design that wove together contribution analysis, realist evaluation (reading each case as a context, a mechanism and an outcome), thematic analysis of the interviews, a partner survey and a large body of programme documents, at country, regional and global levels.

What we look at here is one strand of that work: how the Causal Map team supported Southern Hemisphere to bring causal mapping into the evaluation. Causal mapping means collecting, coding and visualising the interconnected causal claims, the "links" between "factors", that people make in narrative data like reports and interviews. Each link records one claim that someone said one thing influenced another, and it keeps the verbatim quote behind it. Assemble enough of them and a map appears: a picture of the routes along which people said change travelled, from the programme outwards. Think of it as charting the river of impact, with all its tributaries and backwaters, rather than drawing the canal someone planned. And because the evidence here ran to many thousands of claims, the Causal Map team did it with AI for scale.

This is a joint account, written from both sides of that collaboration. When we say "we" we mean the three of us looking back together; where the difference matters we say "the Causal Map team" for the mapping and "Southern Hemisphere" for the evaluation. And we enjoyed the work. The evaluators pushed hard on the method, which is exactly what you want, and it held up.

What do you do with too much evidence?

The usual complaint in evaluation is that there is too little evidence. Here we had the opposite, which is the nicer problem to have, until you sit down to write the report. Coding all the narrative data produced over 22,000 causal claims; even after setting aside the planned and hypothetical ones and keeping only claims about what factually happened, 13,756 links remained to make sense of. A team reading and hand-coding at that scale would be exhausted long before reaching a defensible overview, and human coders get tired, and bored, and pay uneven attention. You could sample, but then you are back to guessing which stories matter.

So the Causal Map app uses a large language model as a low-level coding assistant rather than an oracle. It does the patient work of finding and coding every causal claim across all 176 documents, while the analysts write the coding instructions, check the output against the quotes, and decide what the map means. The AI is tireless where humans tire, and every claim it codes can be traced back to the sentence it came from. That brings the real question into focus. It is no longer how to squeeze a conclusion out of too little. It is how to keep faith with too much: how to hold thousands of separate voices in one view without flattening them into a slogan or cherry-picking the quotes that flatter the programme.

The fair case for the alternatives

Before claiming causal mapping helps, it is worth stating plainly what it is up against, at its strongest.

The **linear theory of change** in Figure 1 is not a foolish thing. It communicates intent at a glance, it disciplines planning, and funders rightly ask for one. Its limit is that it assumes value sits only at the top of the ladder, and it has nowhere to put the surprising, sideways and circular things that real programmes do. A movement-building programme is mostly tributaries and feedback, and a ladder cannot draw a loop.

Traditional qualitative analysis reads the transcripts closely and respects their texture. It is the gold standard for depth. But it does not scale to thirteen thousand claims, and when an evaluator summarises that much material by hand, no reader can check which quotes the summary rests on.

And there is the fair objection that **a diffuse movement across ten countries is simply too complex to evaluate**. That worry deserves respect. The answer is not to pretend otherwise. It is to be modest about what the method claims: causal mapping does not measure how much difference the programme made. It organises the evidence so that a human evaluator can see the pathways and weigh them. That is a smaller claim, and a more defensible one.

What the Causal Map team delivered, and how the report used it

Our contribution started with that coding pass over the whole corpus. From its results we built maps at three scales, one for each of the ten countries, one for each region, and one global map, so a finding could be stated nationally, regionally or programme-wide. We cut the same data thematically too, into maps for capacity strengthening, movement building, advocacy, policy influence and sustainability, each feeding the matching chapter's "Causal map analysis". Every map came with a written commentary tracing its strongest pathways, counting the links, and saying plainly where the evidence was solid and where it was thin. We compared the countries against one another, and found that nine factors were mentioned markedly differently from one place to another: "key populations have increased capacity" stood out in Mozambique, while "a hostile legal and political environment" loomed large in Uganda and barely registered in Egypt. We traced the feedback loops, the virtuous circles and the vicious ones discussed below. And we set the programme's official theory of change against the pathways people actually described, which fed the report's verdict on whether each of its assumptions had held.

In the report's words, the maps were used "to visualise and test causal dynamics across large volumes of qualitative and secondary evidence and to support programme-wide sense-making", in collaboration with the Causal Map team, and they fed the validation workshops where findings were checked with the task force and in country.

The maps were also only one of three lenses. The evaluators set the macro picture from the causal map against the micro picture from each country's contribution story (its context-mechanism-outcome reading), and against a thematic analysis of the same interviews. As the report puts it, change in Love Alliance was "analysed in three different ways, and on different scales, but using mostly the same data sources." Where the three agreed, a finding was safe; where they diverged, that was worth chasing.

Figure 2 is the overall map drawn from the country data. The size of each box and the numbers on it show how often that factor was mentioned; the arrows show claimed influence; darker green factors sit nearer the ends of chains (outcomes), lighter ones are drivers; arrowhead colour shows whether the claimed effect was positive or negative.




Figure 2. An overall causal map built from what people said. The largest box, "Love Alliance provides support for advocacy and capacity building", is the most frequently mentioned factor and the dominant driver. Numbers are citation counts.

Read it and a pattern that no ladder would have predicted comes through. The programme's support flows first into local partner organisations, which build the capacity of the communities they serve; that capacity then feeds advocacy, peer support and a stronger movement. The dominant story practitioners told was rarely "we delivered outputs and outcomes followed". It ran closer to "we built each other up, and that was the point". The maps put a number on the backbone of that story: across the whole dataset the single most-cited factor, with 777 citations, was Love Alliance support for advocacy and capacity building, from which two "workhorse" pathways ran, capacity-strengthening and networking. The report built its central conclusion on exactly that.

One pathway shows how the counts fed the findings. The advocacy map showed a strong connection from Love Alliance support to partners running advocacy campaigns, cited 44 times, leading on to reduced stigma and discrimination against key populations (37) and to key populations empowered to advocate for their own rights (29), with spillover into partner capacity (32) and community training (31). The report carried those counts verbatim into its effectiveness chapter, using them to argue that partner-led advocacy was a central driver of change, enabled by Love Alliance funding and by the networking around it.

The maps also showed something the planners' ladder cannot: reinforcing loops. In the movement-building view (Figure 3), networking and collaboration sit at the centre as a hub, and two virtuous circles appear, one where growing community capacity feeds peer support which feeds capacity again, and one where networking enables advocacy campaigns which generate more networking.




Figure 3. The movement-building view. "Partner organisations implement networking activities, collaborations and build alliances" is the central hub, with reinforcing feedback loops around it.

One respondent, in a group interview in Zimbabwe, put the hub plainly: "The coalitions and partnerships also worked, there is indeed power in numbers ... It was actually this strong partnership which also made things to work for us."

The dataset is overwhelmingly positive in tone, which is worth seeing rather than assuming. Of the factual links, far more were coded positive than negative.

What was coded	Count
Documents used (of 193 coded)	176: 79 country reports, 97 transcripts
Causal claims coded in total	over 22,000
Factual links used in the maps	13,756
... positive / negative / neutral	10,366 / 2,960 / 151

Reading a map source by source

A map of thirteen thousand links is only the beginning of the work. The real work is interrogating it: for a given pathway, which sources actually told that story, and what did they say? The Causal Map app supports this directly. You can write a vignette, a short commentary on a chosen pathway from the perspective of the individual sources behind it, and you can trace a path back to the sources and the verbatim quotes that evidence each step. This is where a map is tested and enriched rather than just admired, and it is much of what we contributed beyond the maps themselves.

Take a real example. Partway through, a fair worry surfaced: was Love Alliance's advocacy programming itself provoking a backlash against key populations? Rather than argue from impression, we put the question to the data and traced it through the relevant sources, 43 of the 116 in that filtered subset. The evidence pointed the other way, and did so uniformly. In every account that mentioned backlash or anti-rights movements, Love Alliance's support appeared as a force helping communities resist it; none portrayed it as the cause. A partner connected to the Dutch Embassy in South Africa described the support as "also helping to counter the backlash against women's rights and LGBT rights" (source 000121-FRICAD). A focus group said it ensured "that communities are empowered to resist and respond to anti-rights movements in meaningful ways" (000090-RIESGA). In Kenya, strengthened legal awareness was "especially vital in the face of growing anti-LGBTIQ movements ... where communities have used their strengthened legal awareness to challenge discrimination and defend themselves more effectively" (000093-NAPVIR). In Nigeria, a participant valued help to "understand the legal framework that is Prohibiting ... and then how do we push for some kind of incremental change" (000107-HISSIN). In East Africa, the support "has helped ensure that advocates have safety nets, can challenge criminalization laws, and can educate policymakers" (000073-EASTAF).

That does not mean the maps painted everything rosy. Filtering the same data for negative claims surfaced a real and important dynamic: a hostile legal and political environment that hinders advocacy and, worse, reinforces itself. The report drew this straight from the map as its primary threat to sustainability, a self-reinforcing negative loop with a count of 232 and an average sentiment of -1.00. The maps also caught the harder truth that empowered advocacy can itself provoke more hostility, which then dampens advocacy: a backlash loop visible only once you filter for negative-sentiment links. The method could hold both the positive pathways and the real risks in the same evidence base, and tell them apart by going back to the sources.

What causal mapping does not do

Stating the limits is part of the method.

The maps are heavily filtered. To stay legible they show only the most-mentioned factors and links, so the absence of a link on a map does not prove the absence of that influence in the world, only that it was not among the most cited. A citation count measures how often something was said. It says nothing about how strong the effect was: a pathway mentioned a hundred times is not therefore twice as powerful as one mentioned fifty times. And longer chains, from support all the way to reduced stigma, are always less cited than the first short steps, so the downstream outcomes rest on a thinner evidence base. Say so, every time.

Above all, this is not causal inference. The map does not turn "twenty people said so" into "therefore it is so". It identifies and organises the evidence; the warranting stays the evaluator's job, and is best done with an explicit, written discipline (see [the quality-assurance companion](#) and [the minimalist coding stance](#)). Reading a map as if it already proved its pathways, especially across sources who each told only part of the story, is [the transitivity trap](#), and it is the single thing most worth guarding against.

Why this helps

For the practitioner, the gain is concrete. You can hold thousands of voices in one picture, at a scale no hand-coding could reach, and still keep faith with each one, because every link on the map carries the quote it came from. You can show a funder the pathways people actually described, loops and all, next to the ladder they planned, and have a frank conversation about the difference. When a doubt surfaces, like the backlash question above, you can put it to the data and trace it through the sources rather than trade impressions. And the maps are live: each one is a bookmark in the Causal Map app, so anyone who doubts a link can click through and read the evidence behind it.

The ladder tells you where the programme hoped to go. The map shows you the country people actually travelled through, and lets you decide, in the open, how much of the journey you believe.

Ready for more

We have tended to describe causal mapping as a preparatory step: we organise the evidence so that someone else can do the evaluation. The companion paper on quality assurance makes the same modest case, and rightly, because the final evaluative judgement belongs to the evaluation team. But look at what the maps did here. They tested the programme's theory of change against the pathways people actually described. They named the main threat to sustainability. They settled the backlash question by going back to the sources. That is evaluative work, done in the open and backed by quotes at every step. So from our side at Causal Map, we know we are ready to take on a bigger part of the evaluation burden: working alongside evaluators rather than standing in for their judgement, but carrying far more of the analysis and sense-making than the words "preparatory step" suggest.

Related

- [the case studies chapter](#)
- [Quality assurance at each step of the causal coding workflow](#): how to warrant conclusions at each step
- [Minimalist coding for causal mapping](#): why we code links simply
- [The transitivity trap](#): the main hazard when reading pathways