

Rubicon principles

Work in progress. The principles Rubicon works to, taken from its design document and kept in step with it. What Rubicon is, and how a piece of work goes, are on rubicon.causalmap.app.

Principles

The rules that apply to writing a workflow, the recipes and the named methods among them, are on the drafting page. What follows applies to everything Rubicon does.

Worth is declared by a person

Rubicon never decides what counts as good. A person sets the standard, in advance and under their name; the AI applies it to the material and shows its working.

Julian King put the objection in October 2025: AI should not be making evaluative judgements about human affairs. The evidence-synthesis bodies draw the line in the same place, where a model makes or suggests a judgement (Flemyng & {Noel-Storr 2025}), and the case that the technology can strengthen qualitative inquiry is being made alongside it (Wise et al. 2026). We agree.

- Looking at the material cannot tell you where adequate begins, and a standard set after seeing the numbers can be made to fit any of them. So the standard comes first, from somebody who put their name to it.
- **Offer the form of a standard, never its numbers.** Ruby does not supply a threshold of her own, such as "overwhelming means eight of ten": she shows where a threshold goes and asks the evaluator for it.
- A registered workflow is versioned and never replaced, and records who registered it and when (on the run path the machine writes this). Freezing a version, and recording whether the freeze came before the run, is designed but not built.

A step is making an evaluative judgement when all three of these are true. They can apply to a coding column ("code the passages where the partnership was meaningful") or a counting rule (a threshold of more than half) as well as a rubric. The app does not test them: it treats a step of type **judge** as the evaluative one, asks for a standard in its rubric, and does not remark on a value declared anywhere else.

- It **restricts** a set of options declared in advance, usually to one of them, sometimes to a band such as "adequate or better".
- Each option **states what it takes to earn it**, in enough detail and context that somebody could hold it against real material and say whether the material meets it; a description that only gestures leaves the reasoning to be rebuilt case by case. "Green: every partner described a change they could date and name a cause for" is a standard; "Green" alone is not. The engine refuses a coding column whose levels are bare names. A rubric that lists its verdicts must state what each stands for; one that lists none runs with the meanings reported as missing.
- The options are **graded by worth**. Health, farming and income is a classification. A frequency table over speculative, reported and evidenced is graded but chooses no point on the distribution.

at_least works on an ordinal and never on a category. The test applies to the answer a step gives, so a workflow is usually descriptive in the middle and evaluative at the end; structurally, that means whether it ends in a judge step.

Transparency is not enough on its own. There has to be a place to disagree: override the verdict, rewrite the rationale, reject a cited quote, mark a criterion unanswerable. The override would be shown beside the machine's verdict, record who and when, and be what everything downstream uses. None of this is built: the only human decision the app records is accepting or withdrawing a whole result.

Every claim points at a quote

- Any narrative or verdict about the material must lead back to its quotes (Moravcsik 2014). The engine refuses to store a write-up or a judgement that cites none, by requiring the coding to be a declared input of the step making the claim. A table is not held to this.
- This is the likeliest way a tool like this fails, and the failure is invisible: two steps from the source the prose still reads well and the numbers still look precise, with nothing connecting either to a document. Note it's easy for an AI tool to reach a conclusion and then search the text for quotes to back that up. Rubicon does not fall into that trap.
- **Some findings are about the method, not the material:** whether contrary evidence was looked for, whether the sample supports the claim. These cite the run record, and the judgement records which kind each verdict was.
- A criterion should be one kind or the other; the engine does not refuse one that is both. "Was anything found that cuts against the hypothesis, and how was it handled?" is two questions.
- A band should stand for one state of affairs; the engine does not read a band's description. If the worst band covers both "contrary evidence was found and it is damaging" and "nobody looked", the reader cannot tell which happened.

A negative finding is the hardest case, and the test is what the claim CITED, not what it was handed. A write-up that cites nothing, such as which of a theory's links the material never mentioned, is backed by the search: what was looked for, over what frame, how much was read, and how many rows it found. This is in the asset's payload as **search_record**, beside the write-up in the export.

Give every proportion its base, and include in the base the documents that gave nothing. A denominator built from coded passages drops the silent documents, so "of the sources that gave me something" reads as "of the sources" and a sparse finding, such as unintended harm, looks bigger than it is.

For a presence question, code the presence. Where a question asks whether something is there, such as whether households keep livestock, code only the passages where it is; a document without a coded row did not mention it. To count the documents that did, give the number a **where** listing the values of a column every coded row fills. Without one, a count over units includes the silent documents; the engine warns on such a number, under "Working with numbers". A level meaning "not mentioned" does not have a passage to quote, so its rows cite unrelated text, and every document gets a row. Keep a third level such as "unclear" only where "the document addresses this and is unclear" is itself a finding.

Beside what it found, the engine reports the passages dropped because their quote could not be located, how much of each document was examined, and how many model calls it made and what they cost.

Build the thing, and not only the verdict

A verdict compresses a whole analysis into one word, and tracing back does not rebuild it. Where a reader would expect a graph, a table of outcomes or a cross-tab, the workflow builds it and declares it as a result; it is usually a step on the way anyway. For a question about whether a theory of change held, that means the table of supporting and contradicting passages per link, beside the verdict. A table can be accepted as the answer as readily as a verdict. A question that wants a number should not pay for a graph nobody asked for.

Where the evaluator does not want criteria settled before the reading, because the material may be richer than the question anticipated, the workflow stops at the table or the prose and reaches no verdict. A free summary ("summarise the findings") is a step away from well defined, since another reader would write another one; where a summary is wanted, specify it narrowly, such as the main point of each finding.

Say what you are comparing against

A programme did well: compared with what? There are seven answers, and they are not interchangeable.

- **A denominator.** Eleven of the fourteen partners. Needs a unit and a stated base.
- **An alternative.** Which of two readings the material better supports.
- **A written standard.** A rubric agreed in advance. The only one that works when there is nothing to compare with.
- **A comparison case.** The programme before it started, a similar one elsewhere, a site that got no support. Depends on the cases being alike in the ways that matter.
- **What was expected.** The theory of change, the target, the plan. Measures delivery, not worth.
- **What would have happened anyway.** Usually the hardest; approachable through causal mapping and other generative accounts of causation. See the method page [generative-and-counterfactual](#) for when each framing fits.
- **The evidence itself.** Sometimes the finding is that nothing here is well enough evidenced for any of the above.

Counting how many people said something does not settle whether it was any good.

Search both sides equally. Iterating on what you have read can fit the method to one corpus: an instruction narrowed until it finds what you saw, a threshold set where the numbers fell. Searching both sides in one pass over the same text is a convention the app does not enforce. Where two codings are compared, it checks they had the same opportunity, per align, and reports any disparity without refusing.

Hold material back. Develop the workflow on a subset, or on analogue material from a previous year or a similar programme, then run the settled workflow over the rest. Where the workflow writes its own tests, split the corpus so that the half that wrote a test never scores it. Before the real material, a workflow can be run over invented documents written to score well and badly (page synthesise-step); if it does not separate them, the method is at fault rather than the material.

Say what you did not read

A workflow may read part of the corpus, by drawing a sample of documents or by narrowing to the passages about a subject; it must not report the result as though everything had been read. A number's frame is whatever reached the step that produced it.

Every result records its frame. The sample step records how many documents it read out of how many, what share of their text it passed on, and what chose that share (a screen, a meaning search, a pattern, a budget). The runner stamps this on every result built on it, through the lineage, and the page and the report print it above the numbers, the verdict and the write-up: a count from part of the text is a floor for the corpus, and a share or a comparison is over what was read. A result whose inputs were read differently lists each frame. In the open format the frame is each result's **coverage**, and a run's budget and the share of text read within it are on its **run**.

What a step loses is part of the frame too. A chunk whose reply could not be read, a row that failed the declared columns, a row a soft recode could not place because its label went under two values: the step adds to its frame how many passages of how many, why, which documents they came from, and, where known, which values they would have counted towards. Later results inherit this. A count of such a value is then a floor, and every step that writes about numbers is told so.

What a reading is shown is part of the frame too. A **judge** or **write** is shown every quote in its inputs, up to a thousand unless its **quotes_shown** sets another limit. Beyond that it is shown a spread, one quote from every unit before a second from any, and told how many of how many it is seeing; the result records both numbers.

State a number inside its frames, innermost first: 28 of the 718 achievements coded mention the embassy, those 718 came from 23% of the text of 142 documents, and a meaning search chose that 23%, so 28 is a floor for the corpus. Whenever the reading was partial, every step that writes about numbers, a judge's rationale and a write-up alike, is given the frame and this form, and Ruby states results the same way in conversation.

Reading a corpus covers when narrowing is right and when it is wrong.

After narrowing, state an absence as an absence in the search: "No passage above 0.45 on childcare, in any of the thirty-seven interviews searched", never "Nobody mentioned childcare". The asset records what was looked for, how close a match was required, and which documents contributed nothing. The method page **reading-enough** covers how much has to be read before an absence or a proportion can be relied on, and how a search misses a link that is in the text.

What it refuses, and what stays soft

- **"Read the documents and tell me whether the project was effective."** In one pass, that gives a confident verdict with no stated standard and nothing specific to disagree with.
- **"Summarise these documents" and "what are the main themes"** are the same fault at a smaller size: a global judgement with no procedure. They get a workflow, a coding in the respondents' words and a clustering, never a direct answer.
- **"Is this a good project?"** is four questions: good for whom, compared with what, against whose standard, and how much of that this material can answer. Where a terms of reference uses the generic criteria (coherence, efficiency, sustainability, relevance, equity), the method page **generic-evaluation** turns each into a workflow with its own standard.
- The test is whether the question is well defined. A question is never refused for being about tone, feelings or anything else that sounds soft.
- **Refusing is rarer than it sounds.** "What percentage of the budget went on training?" is answerable where a report states the number. The usual answer is a workflow that goes and looks; if nobody wrote the number down, that is the finding, backed by the search.
- **Some questions are well defined in parts and not as a whole.** "What share of the fellows who wanted a policy meeting got one?" is two codings and a ratio. A single pass would be arithmetic done in a model's head.
- **A question whose answer is in a spreadsheet is refused.** Hours between landfall and the first cash payment, or cost per person reached, come from a payments system without any words behind them, and every number here must be traceable to words somebody said. The workflow states which part of the question it answers, and you bring the other numbers in when you write.
- **A number spoken in a document is in scope.** "We reached about four hundred children", said by a school director, has a speaker and a position in a source. A coding step declares a numeric column, and the counting does arithmetic on it.
- **A judgement may refuse the question.** Where the material shows the question assumes something that did not happen, or compares things the sources cannot compare fairly, the judge step reports this instead of awarding a band, and cites the evidence.

Decline the statistics Rubicon does not do

Ruby's stance. Rubicon does not compute effect sizes, differences in differences, power calculations, regressions, or any number for what would have happened without the programme; a version improvised from counts would be wrong in a way no reader could see. So she declines, plainly and without apology, names the kind of analysis needed, says it has to be done from the data by somebody with a statistics package, and offers the part Rubicon can answer, usually what people said about the change and why. She does not name a product or a website, or write a weaker version to have something to show.

The house phrase, to adapt rather than recite: "Answering that properly needs [a difference-in-differences estimate, an effect size, a power calculation], worked out from the data by somebody with a statistics package. Rubicon reads and counts what people said, so what it would give you is a number that looks like an answer without being one. What it can do is [the part it can answer]."

Where the refusal depends on a fact about the material, such as there being no comparison group, she says where that fact comes from, since she has read none of the documents: a column the sources have, or the absence of any column marking who was reached. She never writes "Every household was reached by the programme" into a frame description, where it would enter the record as a finding. The form is "nothing in the sources' columns marks a household the programme did not reach, so the material offers no comparison group". A coding step that reads for it is the way to find out.

Where the question is about a comparison, she proposes a number whose groups are a property of the documents, such as who was interviewed, since that can be tested, per Numbers. A difference the test says could easily be chance is not established, however it looks in the table: say so, or leave it out. Nor does it show the groups are alike, since with small samples most real differences would not show either.

Where this sits in the literature

- Scriven's logic of evaluation separates description, "What's so?", from valuing, "So what?", then reaches a verdict through criteria of merit, standards of merit and a synthesis into overall worth (Scriven 1980, 1991; Fournier 1995).
- Davidson: an evaluation must ask not only "What were the results?" but "How good were the results?" Her worked failure is an executive summary of counts and percentages with no verdict in it (Davidson 2014).
- Gargani and King's second principle: "an evaluator cannot ascertain the value of an impact solely by studying it because value is not located within it". The philosophical question is open; Putnam's argument against the fact and value dichotomy is the standing objection.

References

Flemyng, & {Noel-Storr (2025). *Position Statement on Artificial Intelligence (AI) Use in Evidence Synthesis across Cochrane, the Campbell Collaboration, JBI and the Collaboration for Environmental Evidence 2025*. Cochrane Database of Systematic Reviews.

<https://doi.org/10.1002/14651858.ED000178>.

Moravcsik (2014). *Transparency: The Revolution in Qualitative Research*. APSCC Newsl., 47, 48--53.

<https://doi.org/10.1017/S1049096513001789>.

Wise, Gresalfi, & {Spencer-Smith (2026). *Why AI Is Not the Enemy: Opportunities to Strengthen Core Commitments of Qualitative Inquiry Through Trustworthy AI-in-the-Loop Analysis*. International Journal of Qualitative Methods, 25, 16094069261435579. <https://doi.org/10.1177/16094069261435579>.