🌻 Map Panel

🗺️ **What you can do here:** See your causal relationships as an interactive network map. Drag nodes around, click on links to edit them, and use the controls to customize how the map looks. You can even drag one factor onto another to quickly create new links. This is where your data comes to life visually.

Map Controls 🎚

- 👉🏻 **Jump to factor** (type-to-search dropdown): quickly find and select factors on the map (supports multiple selections).
- 👉🏻 🔄 **Refresh layout** (button): reset the map layout after zooming/moving.
- 👉🏻 📷 **Copy image to clipboard** (button): copy a high-quality map image for reports/slides.
- 👉🏻 📋 **Copy legend** (button): copy the map legend text.
- 👉🏻 **Zoom in/out** (controls): zoom the map view.
- 👉🏻 **Double-click** (gesture): zoom in to that point on the background.

Map Legend ☰

Discrete text legend showing:

- projectname and included sources
- Citation coverage percentage
- Visual encoding explanations (link sizes, colors, numbers)
- Applied filters summary
- 💡 Tip: Click Copy legend to copy this text to clipboard.
- You can drag the legend box to reposition it on the map.

Evidence (shortcut to Links Print View)

- 👉🏻 **Evidence** (button, bottom-right on the map, opposite the legend): switches to the **Links** tab and opens **Print View** for the **same filtered links** you see on the map. It applies a preset—**Group by** Bundle and Source, **sort** by citation count (highest first), **Show details** and **Context** on, **page size** 50—and updates the URL so you can bookmark or share that layout. A toast links to Print View help; a short hint may point to the **Print view** toggle.

© Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

Map Formatting 🎚

Customisable formatting (Things you can tweak)

Layout and interaction

- 👉 **Layout** (dropdown): choose how the map is laid out.
- Interactive and most other layouts are good while you are conducting your research (fast + supports the interactive features).
- Print/Graphviz is best for static images (reports/journal articles). In Graphviz SVG you can still pan/zoom (mouse wheel, double-click, Shift+double-click).
- 👉 **Groups** (switch): layout maps with top-level factors as boxes which group together their "children".
- 👉 **Initials** (dropdown): add distinctive coloured badges to factor labels in Interactive and Print layouts. Choose which part of the label is used for the badge: Off / Full label / Level 1 (;) / First colon (:) / Square brackets [] / Round brackets (). Uses the same extraction rules as Groups. In Print layout, badges use a gradient for visual distinctiveness.
- 👉 **Direction** (dropdown): LR (default), TB, or BT (for Interactive and Print/Graphviz layouts).
- 👉 **Link direction** (dropdown): Normal (directed arrows) vs Undirected (dots at both ends).
- In Undirected mode, dots use the same colours as arrowheads (including sentiment colouring). When sentiment is neutral (0), they use **Link Colour**.
- Note: when the Combine Opposites filter is active, tail/head can still have different colours.

Factors

- 👉 **Factor labels** (dropdown): what to show next to each factor (same data as the Factors Panel).
- Citation count (default) / Source count / Sentiment (mean incoming) / None
- 👉 **Factor colours** (dropdown): Outcomeness (default) / Source count / Citation count / None
- 👉 **Factor sizes** (dropdown): Citation count (default) / Source count / None

Links

- 👉 **Link labels** (dropdown): what to show on each link.
- Citation count (default) / Source count / Sentiment / Label by Group / Unique Sources / All Sources / Unique Tags / Unique Tags (Tally) / All Tags / None
- 👉 **Link widths** (dropdown): Citation count (default) / Source count / None
- 👉 **Link label font size** (control): change link label font sizes.

© Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

- 👉 **Arrowhead size** (control): scale arrowhead size (Interactive + Print/Graphviz). Default 100% keeps current appearance.
- 👉 **Link colour** (colour picker): sets the default link line colour (Interactive + Print/Graphviz). When sentiment is neutral (0), this colour is also used for arrowheads and node borders.
- 👉 **Links highlight** (dropdown): optional extra highlighting without changing the base colour scheme.
- Off (default)
- Reverse (backwards/same-rank in current layout direction)
- Significant (when Label by Group shows ⬆/⬇)
- Feedback loop (2 / ≤3 / ≤4 factors)
- Feedback loop + reverse (combine the above)

Other

- 👉 **Show self-loops** (toggle, default on): show/hide A→A links on the map.

Fixed visual appearance (things you can't tweak)

Some parts of the map's appearance are automatic (i.e. they are not controlled by the Map Formatting widgets above):

Link geometry (bundling):

- Links are bundled and drawn as curved edges for readability.

Automatic colouring overlays:

- Arrowhead colours reflect mean **sentiment** for that link bundle (neutral uses your chosen **Link colour**).
- When the Combine Opposites filter is active, arrowhead colours instead reflect **flipped share** (tail=cause, head=effect).

Automatic highlighting:

- Factors that match filters like Factor Label or Path Tracing show dashed coloured borders.
- Factor border colour reflects mean incoming edge sentiment (but when Combine Opposites is active: average flipped share, blue→red).

Interactive Features

These work for all layouts except Print/Graphviz layout (which is mostly for static export, but does support clicking nodes/links now).

© Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

- **Drag factors** to temporarily reposition them
- **Drag factor to factor** to create new links
- **Shift+drag** for box selection of multiple factors (opens edit modal)
- **Ctrl+drag** for box selection of multiple factors (direct selection, no modal)
- **Click a link** to edit.
- **Click a factor** to edit; shift-click or ctrl-click to add to selection without opening modal.

Editing and deleting (multiple) factors

- Select factor(s) by clicking a factor, shift-click or ctrl-click to add more, or shift+drag/ctrl+drag a box around multiple factors, then:
- Move selected factors together
- Delete matching factors everywhere or in current view only
- Rename matching factors everywhere or in current view only

What does "everywhere or in current view only" mean?

- **everywhere**: all links containing factors with exactly the selected labels will be deleted
- **in current view only**: all links containing factors with exactly the selected labels (and matching the current filters, i.e. those you can see in the current map) will be deleted

💡 Tip: By control-clicking or shift-clicking multiple factors you can easily rename several at once, e.g. you can merge multiple factors as a single factor.

Grid layout

Factors containing a tag of the form `(N.M)` or `(N,M)` anywhere in the label (where N and M are integers) are positioned on a grid layout. The grid coordinate tags are automatically stripped from displayed labels. Grid tags can also be **partial**: `(N,)`, `(,M)`, `[N,M]`, `[N,]`, `[,M]` (same meaning; first number = rank direction, second = perpendicular).

Grid layout toggle: Enable/disable grid layout in Map Formatting. Defaults to enabled. Disabled automatically when no grid tags are present.

Interactive Layout:

- Grid-tagged factors are positioned at their grid coordinates and locked in place
- Other factors with no grid tag are positioned freely within the grid bounds
- Grid bounds: from smallest x -1 to largest x +1, and smallest y -1 to largest y +1

Print/Graphviz Layout:

- Grid-tagged factors anchor the initial and final ranks:

 © Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

- Factors with minimum rank coordinate (first number) are anchored at `rank=min` (initial rank)
- Factors with maximum rank coordinate are anchored at `rank=max` (final rank)
- This improves layout stability while allowing Graphviz to position other nodes optimally
- Grid coordinate tags are stripped from labels in the output
- The **perpendicular coordinate** (second number) is not an absolute y-position in Graphviz; it is only used as a **best-effort ordering hint within a fixed rank** (so y-only tags like `(,M)` cannot be enforced unless the rank coordinate is also specified).

Grid coordinates respect layout direction:

- **First number (N)** always maps to the rank direction (main flow direction)
- **Second number (M)** always maps to the perpendicular direction
- **BT (Bottom-Top)**: First number = y (rank), y starts at bottom (flip y), second = x
- **TB (Top-Bottom)**: First number = y (rank), y starts at top (normal), second = x
- **LR (Left-Right)**: First number = x (rank), x starts at left (normal), second = y, y starts at top
- **RL (Right-Left)**: First number = x (rank), x starts at right (flip x), second = y, y starts at top

Vignettes 🖊

📝 **What you can do here:** Generate AI-powered narrative summaries of your causal maps. Choose between a **whole-map** summary that covers the relationships in your current filtered view, or a **typical-source** story that focuses on one representative respondent. Useful for reports, annexes, and explaining your analysis in plain language.

For practitioners (getting started)

What are vignettes and why use them?
Vignettes turn your coded causal map into readable prose. After you have links (and usually quotes on those links), the AI can draft a narrative that names themes, tensions, and patterns in your material. You stay in control: you choose the prompt (tone, audience, structure), and you can edit the text afterwards. They are a fast way to move from “many links on a map” to “something I can paste into a report or share with stakeholders”—not a substitute for your judgement, but a structured first draft.

Whole-map vignette
Use this when you want a bird’s-eye story of **everything that is currently in the map** (respecting your filters and source selection). The app sends a compact summary of factors and bundled links (with sentiment), **evidence snippets** (quotes and source IDs from highlighted bundles where available), and up to **30 “typical” sources** scored by how much they represent common bundles—each with bundles, link-level quotes, and a small **metadata preview** (title, filename, and simple custom fields). If you want the narrative to emphasise particular edges, set

 © Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

Map Formatting → Links highlight to **Significant** or **Feedback loop** first; that snapshot is included so the model can focus there. **AI region** (where Gemini runs) follows **Project Details → AI region**, not a separate control on this card.

Typical-source vignette
Use this when you want a **single-respondent case study**: the app picks the most "representative" source for the **current** map view (using link counts and coverage of bundles, with a bias so very long documents do not always win). It sends that source's **full text** plus its links with quotes and sentiment, so the model can write as if telling **one** story. Do not ask it to generalise across the whole project in this mode—the prompt and data are scoped to that one source.

How to use (UI):

1. Open the **Vignettes** card on the Map side panel and expand it if needed.
2. Choose **model** (and optional **thinking** controls if your model supports them). **Region** for AI is set per project under **Edit project** (AI region), not on this card.
3. (Optional) Leave **Enable checking (second AI pass)** on so a checker can review the draft; notes appear in a collapsed panel when relevant.
4. Edit the **Whole-map prompt** and/or **Typical-source prompt** (prompt history: dropdown and prev/next; see tips on using prompt history).
5. Click **Write Whole-map Vignette** or **Write Typical-source Vignette** (two separate buttons).

Tip (optional): tell the AI which links matter most (whole-map mode):

- Go to **Map Formatting → Links highlight** and choose **Significant** or **Feedback loop**.
- When you run the **whole-map** vignette, the app includes a small snapshot of those highlighted edges so the model can focus on them.
- If **Links highlight** is **Off** (or **Reverse** only, with no edges flagged), or nothing qualifies, **no extra highlight list is sent**.

What each mode sends (summary):

- **Whole map:** Aggregated nodes/edges for the current view, optional **links highlight** snapshot, **bundle evidence** (quotes/source IDs sampled from top bundles so large projects stay within limits), and up to **30 typical sources** with bundles, per-link quotes, and compact source metadata—not the entire raw links table.
- **Typical source:** One chosen source's **full text**, its links with quotes and sentiment, simple node frequencies, and optional links-highlight context for that source.

 © Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0

Output format: Markdown (headers, bold, lists, blockquotes, code blocks). You can copy from the result area.

Bookmarking & restore:

- Each run saves a **bookmark** for the current view (description: `Vignette (whole|typical): <your prompt>`) and appends a link at the bottom of the vignette output.
- The bookmark footer includes the **model name** used.
- Vignette-related settings are stored in **URL state** where applicable (model, thinking, checking, prompts); project **AI region** lives in project metadata.

© Causal Map Ltd 2026 · causalmap.app · CC BY-NC 4.0